

# ann-model-1

September 12, 2024

## 1 By prisca

```
[3]: # this bank have an issue of customers just leaving their bank and they want be
      ↪able to identify those customer even before
      #the leave.
      # this project is about building a model to predict customers who are likely to
      ↪leave than bank and vice visa with following
      # features, using Artificial Neutral Network
```

```
[97]: import pandas as pd
import sklearn
import matplotlib.pyplot as plt
import numpy
import tensorflow as tf
import seaborn as sns
```

```
[160]: import warnings
warnings.filterwarnings('ignore')
```

```
[98]: df=pd.read_csv(r'C:\Users\USER\Documents\Churn_Modelling.csv')
```

```
[99]: df.head()
```

```
[99]:  RowNumber  CustomerId  Surname  CreditScore  Geography  Gender  Age  \
0         1    15634602  Hargrave         619     France  Female  42
1         2    15647311    Hill         608      Spain  Female  41
2         3    15619304    Onio         502     France  Female  42
3         4    15701354    Boni         699     France  Female  39
4         5    15737888  Mitchell         850      Spain  Female  43
```

```
      Tenure  Balance  NumOfProducts  HasCrCard  IsActiveMember  \
0         2      0.00             1           1              1
1         1  83807.86             1           0              1
2         8 159660.80             3           1              0
3         1       0.00             2           0              0
4         2 125510.82             1           1              1
```

	EstimatedSalary	Exited
0	101348.88	1
1	112542.58	0
2	113931.57	1
3	93826.63	0
4	79084.10	0

```
[100]: df.shape
```

```
[100]: (10000, 14)
```

```
[101]: dfr=df.drop(['RowNumber', 'CustomerId', 'Surname'],
↳ 'CreditScore'],axis=1,inplace=True)
```

```
[102]: df.head()
```

```
[102]: Geography Gender Age Tenure Balance NumOfProducts HasCrCard \
0 France Female 42 2 0.00 1 1
1 Spain Female 41 1 83807.86 1 0
2 France Female 42 8 159660.80 3 1
3 France Female 39 1 0.00 2 0
4 Spain Female 43 2 125510.82 1 1
```

	IsActiveMember	EstimatedSalary	Exited
0	1	101348.88	1
1	1	112542.58	0
2	0	113931.57	1
3	0	93826.63	0
4	1	79084.10	0

```
[148]: # visualization
```

```
[103]: df.isnull().sum()
```

```
[103]: Geography      0
Gender              0
Age                0
Tenure             0
Balance            0
NumOfProducts      0
HasCrCard          0
IsActiveMember     0
EstimatedSalary    0
Exited             0
dtype: int64
```

```
[104]: df.duplicated().sum()
```

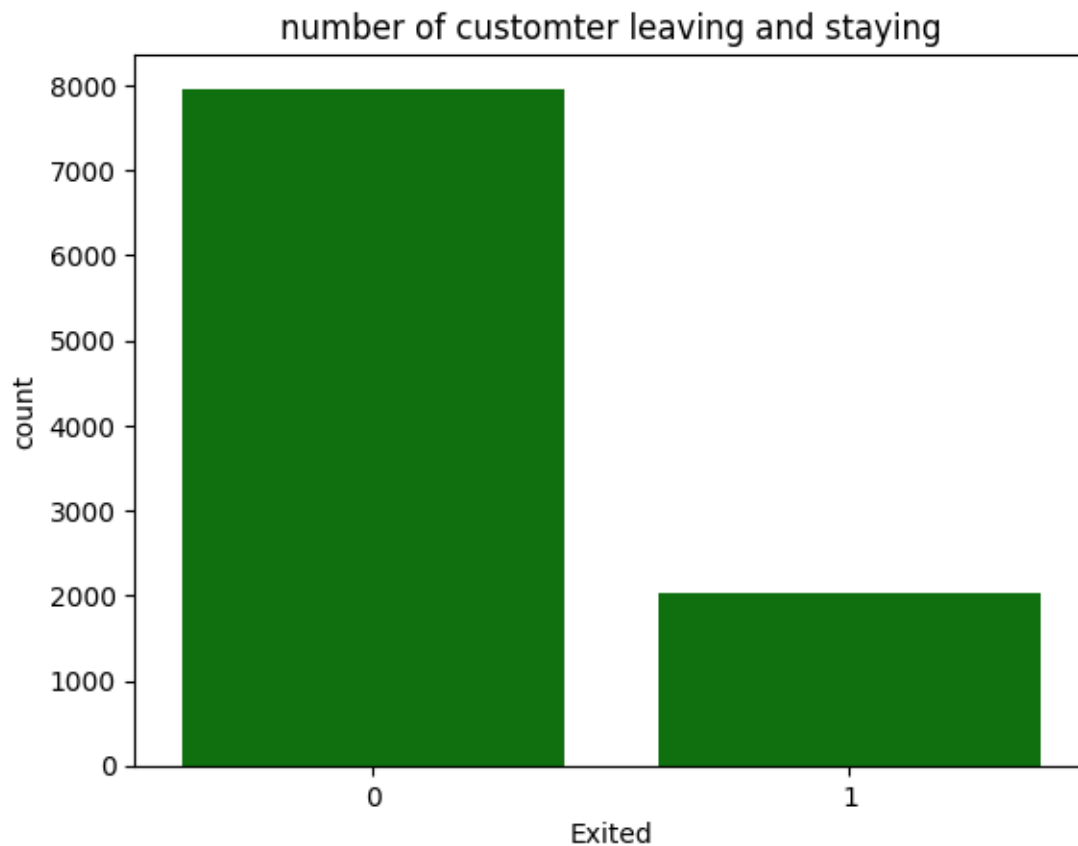
```
[104]: 0
```

## 2 visualization

```
[176]: sns.countplot(x='Exited',data=df,color='green')
plt.title('number of customter leaving and staying')

# the number of customer leaving is less than the number of cutsomer staying
```


```
[176]: Text(0.5, 1.0, 'number of customter leaving and staying')
```



```
[155]: sns.countplot(x='Gender',data=df)
plt.title('male vs female')

# the bank has more male customers than female with little differnce
```


```
[155]: Text(0.5, 1.0, 'male vs female')
```



```
[162]: fig,ax=plt.subplots(figsize=(15,10))
sns.countplot(x='Age',data=df)
plt.title('Age range of the customers')

# the customers ranges from 18 to 92 years old
# the have high amount of customers from age 29 to 47 years
```

```
[162]: Text(0.5, 1.0, 'Age range of the customers')
```

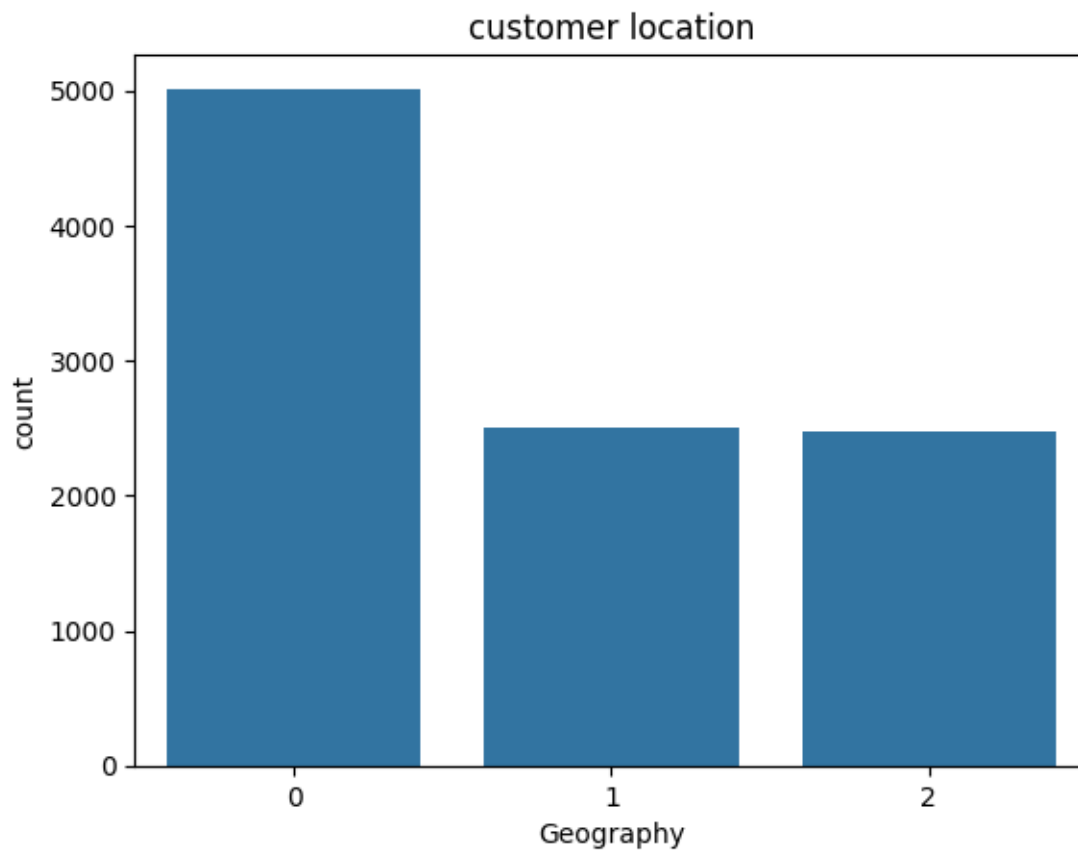


```
[164]: df['Geography'].value_counts()
```

```
[164]: Geography
0      5014
1      2509
2      2477
Name: count, dtype: int64
```


```
[166]: sns.countplot(x='Geography',data=df)
plt.title('customer location')
```

```
[166]: Text(0.5, 1.0, 'customer location')
```



```
[173]: sns.regplot(x='Balance',y='Geography',data=df)
```

```
[173]: <Axes: xlabel='Balance', ylabel='Geography'>
```




```
[169]: df_corr=df.corr()
```

```
[175]: fig,ax=plt.subplots(figsize=(15,10))
sns.heatmap(df_corr,annot=True)

# from the map, there is no high correlation among the variabes
```

```
[175]: <Axes: >
```



```
[ ]:
```

```
[ ]:
```

```
[ ]:
```

```
[105]: from sklearn.preprocessing import LabelEncoder
```

```
[106]: le=LabelEncoder()
```

```
[107]: df['Geography']=le.fit_transform(df['Geography'])
```

```
[108]: df['Gender']=le.fit_transform(df['Gender'])
```

```
[109]: df.head()
```

```
[109]:
```

	Geography	Gender	Age	Tenure	Balance	NumOfProducts	HasCrCard	\
0	0	0	42	2	0.00	1	1	
1	2	0	41	1	83807.86	1	0	
2	0	0	42	8	159660.80	3	1	



3	0	0	39	1	0.00	2	0
4	2	0	43	2	125510.82	1	1

	IsActiveMember	EstimatedSalary	Exited
0	1	101348.88	1
1	1	112542.58	0
2	0	113931.57	1
3	0	93826.63	0
4	1	79084.10	0

```
[110]: from sklearn.preprocessing import StandardScaler
```

```
[111]: from sklearn.model_selection import train_test_split
```

```
[112]: df.columns
```

```
[112]: Index(['Geography', 'Gender', 'Age', 'Tenure', 'Balance', 'NumOfProducts',
          'HasCrCard', 'IsActiveMember', 'EstimatedSalary', 'Exited'],
          dtype='object')
```

```
[113]: x=df[['Geography', 'Gender', 'Age', 'Tenure', 'Balance', 'NumOfProducts',
          'HasCrCard', 'IsActiveMember', 'EstimatedSalary']]
        y=df['Exited']
```

```
[114]: x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.3,random_state=0)
```

```
[115]: scaler=StandardScaler()
```

```
[116]: x_train=scaler.fit_transform(x_train)
        x_test=scaler.transform(x_test)
```

```
[117]: x_train.shape,x_test.shape
```

```
[117]: ((7000, 9), (3000, 9))
```

```
[118]: #building the machine brain
```

```
[119]: ann=tf.keras.models.Sequential()
```

```
[120]: ann.add(tf.keras.layers.Dense(units=6,activation='relu'))
```

```
[121]: ann.add(tf.keras.layers.Dense(units=6,activation='relu'))
```

```
[122]: ann.add(tf.keras.layers.Dense(units=1,activation='sigmoid'))
```

```
[123]: #compile the ann
```

```
[124]: ann.compile(optimizer='adam',loss='binary_crossentropy',metrics='accuracy')
```

```
[125]: #train the model
```

```
[126]: ann.fit(x_train,y_train,batch_size=32,epochs=50)
```

```
Epoch 1/50
219/219 [=====] - 2s 3ms/step - loss: 0.5421 -
accuracy: 0.7667
Epoch 2/50
219/219 [=====] - 1s 3ms/step - loss: 0.4708 -
accuracy: 0.8014
Epoch 3/50
219/219 [=====] - 1s 3ms/step - loss: 0.4451 -
accuracy: 0.8083
Epoch 4/50
219/219 [=====] - 1s 3ms/step - loss: 0.4293 -
accuracy: 0.8164
Epoch 5/50
219/219 [=====] - 1s 3ms/step - loss: 0.4172 -
accuracy: 0.8227
Epoch 6/50
219/219 [=====] - 1s 3ms/step - loss: 0.4048 -
accuracy: 0.8326
Epoch 7/50
219/219 [=====] - 1s 3ms/step - loss: 0.3938 -
accuracy: 0.8401
Epoch 8/50
219/219 [=====] - 1s 3ms/step - loss: 0.3847 -
accuracy: 0.8437
Epoch 9/50
219/219 [=====] - 1s 3ms/step - loss: 0.3777 -
accuracy: 0.8490
Epoch 10/50
219/219 [=====] - 1s 3ms/step - loss: 0.3727 -
accuracy: 0.8507
Epoch 11/50
219/219 [=====] - 1s 3ms/step - loss: 0.3692 -
accuracy: 0.8517
Epoch 12/50
219/219 [=====] - 1s 3ms/step - loss: 0.3666 -
accuracy: 0.8521
Epoch 13/50
219/219 [=====] - 1s 3ms/step - loss: 0.3642 -
accuracy: 0.8529
Epoch 14/50
219/219 [=====] - 1s 3ms/step - loss: 0.3629 -
```

accuracy: 0.8540  
Epoch 15/50  
219/219 [=====] - 1s 3ms/step - loss: 0.3613 -  
accuracy: 0.8537  
Epoch 16/50  
219/219 [=====] - 1s 3ms/step - loss: 0.3600 -  
accuracy: 0.8547  
Epoch 17/50  
219/219 [=====] - 1s 3ms/step - loss: 0.3594 -  
accuracy: 0.8540  
Epoch 18/50  
219/219 [=====] - 1s 3ms/step - loss: 0.3583 -  
accuracy: 0.8559  
Epoch 19/50  
219/219 [=====] - 1s 3ms/step - loss: 0.3574 -  
accuracy: 0.8547  
Epoch 20/50  
219/219 [=====] - 1s 3ms/step - loss: 0.3566 -  
accuracy: 0.8556  
Epoch 21/50  
219/219 [=====] - 1s 3ms/step - loss: 0.3559 -  
accuracy: 0.8561  
Epoch 22/50  
219/219 [=====] - 1s 3ms/step - loss: 0.3550 -  
accuracy: 0.8551  
Epoch 23/50  
219/219 [=====] - 1s 3ms/step - loss: 0.3545 -  
accuracy: 0.8559  
Epoch 24/50  
219/219 [=====] - 1s 3ms/step - loss: 0.3539 -  
accuracy: 0.8560  
Epoch 25/50  
219/219 [=====] - 1s 3ms/step - loss: 0.3531 -  
accuracy: 0.8563  
Epoch 26/50  
219/219 [=====] - 1s 3ms/step - loss: 0.3528 -  
accuracy: 0.8574  
Epoch 27/50  
219/219 [=====] - 1s 3ms/step - loss: 0.3523 -  
accuracy: 0.8584  
Epoch 28/50  
219/219 [=====] - 1s 3ms/step - loss: 0.3523 -  
accuracy: 0.8557  
Epoch 29/50  
219/219 [=====] - 1s 3ms/step - loss: 0.3515 -  
accuracy: 0.8570  
Epoch 30/50  
219/219 [=====] - 1s 3ms/step - loss: 0.3513 -

accuracy: 0.8573  
Epoch 31/50  
219/219 [=====] - 1s 3ms/step - loss: 0.3513 -  
accuracy: 0.8569  
Epoch 32/50  
219/219 [=====] - 1s 3ms/step - loss: 0.3507 -  
accuracy: 0.8576  
Epoch 33/50  
219/219 [=====] - 1s 3ms/step - loss: 0.3504 -  
accuracy: 0.8571  
Epoch 34/50  
219/219 [=====] - 1s 3ms/step - loss: 0.3504 -  
accuracy: 0.8577  
Epoch 35/50  
219/219 [=====] - 1s 3ms/step - loss: 0.3500 -  
accuracy: 0.8573  
Epoch 36/50  
219/219 [=====] - 1s 3ms/step - loss: 0.3497 -  
accuracy: 0.8584  
Epoch 37/50  
219/219 [=====] - 1s 3ms/step - loss: 0.3490 -  
accuracy: 0.8586  
Epoch 38/50  
219/219 [=====] - 1s 3ms/step - loss: 0.3490 -  
accuracy: 0.8583  
Epoch 39/50  
219/219 [=====] - 1s 3ms/step - loss: 0.3487 -  
accuracy: 0.8587  
Epoch 40/50  
219/219 [=====] - 1s 3ms/step - loss: 0.3485 -  
accuracy: 0.8581  
Epoch 41/50  
219/219 [=====] - 1s 3ms/step - loss: 0.3483 -  
accuracy: 0.8576  
Epoch 42/50  
219/219 [=====] - 1s 3ms/step - loss: 0.3481 -  
accuracy: 0.8573  
Epoch 43/50  
219/219 [=====] - 1s 3ms/step - loss: 0.3477 -  
accuracy: 0.8583  
Epoch 44/50  
219/219 [=====] - 1s 3ms/step - loss: 0.3474 -  
accuracy: 0.8581  
Epoch 45/50  
219/219 [=====] - 1s 3ms/step - loss: 0.3474 -  
accuracy: 0.8580  
Epoch 46/50  
219/219 [=====] - 1s 3ms/step - loss: 0.3471 -

```

accuracy: 0.8577
Epoch 47/50
219/219 [=====] - 1s 3ms/step - loss: 0.3469 -
accuracy: 0.8603
Epoch 48/50
219/219 [=====] - 1s 3ms/step - loss: 0.3467 -
accuracy: 0.8600
Epoch 49/50
219/219 [=====] - 1s 3ms/step - loss: 0.3469 -
accuracy: 0.8574
Epoch 50/50
219/219 [=====] - 1s 3ms/step - loss: 0.3463 -
accuracy: 0.8593

```

```
[126]: <keras.src.callbacks.History at 0x223ca10d880>
```

```
[127]: #predictions
```

```
[128]: ypred=ann.predict(x_test)
```

```
94/94 [=====] - 0s 2ms/step
```

```
[129]: ypred=(ypred>0.5)
```

```
[130]: ypred
```

```
[130]: array([[False],
           [False],
           [False],
           ...,
           [False],
           [False],
           [False]])
```

```
[131]: y_pred= [1 if pred > 0.5 else 0 for pred in ypred]
```


```
[137]: #evaluating the model accuracy
```

```
[132]: from sklearn import metrics
       from sklearn.metrics import confusion_matrix
```

```
[133]: cm=(confusion_matrix(y_test,ypred))
```

```
[134]: sns.heatmap(cm,annot=True)
```

```
[134]: <Axes: >
```



```
[135]: cm
```

```
[135]: array([[2290,  89],
            [ 335, 286]], dtype=int64)
```

```
[136]: metrics.accuracy_score(y_pred,y_test)
```

```
[136]: 0.8586666666666667
```

### 3 testing the accuracy

```
[177]: ann.predict(scaler.transform([[0,1,40,3,60000,2,1,1,50000]])) # nicely predicted
```

```
1/1 [=====] - 0s 38ms/step
```

```
[177]: array([[0.03612975]], dtype=float32)
```

```
[178]: ann.predict(scaler.transform([[0,      0,      42,      8,      159660.
    ↪80,      3      ,1      ,0      ,113931.57      ]])) # correctly predicted
```

```
1/1 [=====] - 0s 38ms/step
```

```
[178]: array([[0.9962028]], dtype=float32)
```

```
[147]: # 0.9965 means customer will leave the bank
```

```
[4]: # with this model the bank can detect customers who are about to leave the bank_
      ↪ and take measures to
      #retain them
```

```
[ ]:
```